

# ReverseKLD Efficacy for Small Language Models

October 25, 2024

## 0.1 Abstract

This notebook is a research project that I conducted to test the efficacy of the results of [2] in their 2024 MiniLLM paper with what I will call “MicoLMs” (models with less than 200B parameters).

I first trained an entire 124M with inspiration taken from Andrej Karpathy and GPT-2 as well as Olmo. With this I created a modular LM trained on the 10B token subset of finewebEDU from Hugging Face. As you will see down below, the model outperformed GPT-2 124M on the hellaswag eval with far fewer training tokens. I then took that pretrained model and created a training framework for the MiniLLM implementation of reverse KLD loss with a teacher model to try to get the Olmo 1B model to transfer some of its intelligence down to the MicroLM.

To keep variables to a minimum, the only difference between the pretraining and the rKLD was the implementation of the KLD itself. No architectural changes were made to the model.

Both models were evaluated during and after training on zero-shot and few-shot hellaswag, snli, and GLUE benchmarks to evaluate performance. Comparing the loss in this scenario is not relevant as the loss calculations are completely different.

## 0.2 Student Model Training Details

The initial student model was trained two times. Both models used the ~10B token subset of FineWeb-Edu. The config for the initial model used RoPE embeddings as well as swish activation functions and RMSnorm. The second model switched those out for absolute learning embeddings, gelu, and layer norm. The second version of the training is the bog standard approach to training LLMs and the data proves why. As you can see in the figures, all aspects: training loss, validation loss, and hellaswag accuracy were all better with the second model.

In hindsight, I believe changing RMSnorm to layer norm was unnecessary as it has been shown to keep similar performance for higher levels of efficiency [1]. As for the other changes, I experimented with Swish, GELU, and SwiGLU but ultimately landed on GELU for faster convergence times [3] and lack of added parameters as seen in SwiGLU. The largest factor towards the reduced performance I believe was RoPE embeddings. While RoPE has been very useful in recent language models and meshes well with flash-attention methods, the generalized nature of the method likely required more data to get to the same place as simple learned positional embeddings. Though, RoPE probably has a higher ceiling and has more utility than just capped size inputs.

## 0.3 Hella Swag Eval

The hellaswag eval methodology was based off Andrej Karpathy’s implementation. Because the models are so small, they don’t have the reasoning ability to be shown a list of choices and pick a result

so instead probability comparisons are used. The model is independently called four times, one for each multiple choice answer and the probabilities the model assigns to each choice is evaluated to confirm what result the model “chose”.

Model	HellaSwag
Olmo 7B	0.7338
Olmo 1B	0.6071
GPT-2 1.5B	0.4839
GPT-2 124M	0.2955
Qwen2.5-1.5B	0.6469

These results are what motivate the idea for this test. These larger models have a significantly higher hellaswag scores, especially given the fact that a random result would be  $\sim 0.25$ . The work of Yuxian et al. shows that for larger models (1B-7B) getting knowledge distilled into them is quite fruitful and their methodology of reverse KLD is good at doing so. I will test if these results hold for even smaller models (124M).

The plot shows that the control model, using the given improved training methods and data outperforms the 124M checkmark of GPT-2.

The training process for reverseKLD was computationally expensive as you would have to evaluate all data twice, once through a much larger model as well so the training process was much slower. The model training method was tested with a variety of hyperparameters (length, alpha, lr) but all changes had similar results. The model did not generalize well and actually got worse.

I believe these results occur from the size of the model just being too small. I believe the idea of superposition may have a lot to do with this. It is believed via the idea of superposition that the amount of facts that can be stored goes up exponentially with parameter size because the amount of “directions” that can be linked to characteristics rises exponentially with dimension. Even with the high level of smoothing, the smaller model could not generalize the results in any meaningful way because the larger training models have too great an amount of  $j$  for such a small model. As you can see from (Yuxian et. al) For larger student models, they have enough dimensionality to still distill some of the information from the teacher models.

### 0.3.1 Citation:

- [1] Biao Zhang, & Rico Sennrich. (2019). Root Mean Square Layer Normalization.
- [2] Yuxian Gu, undefined., et al, “MiniLLM: Knowledge Distillation of Large Language Models,” 2024.
- [3] J. Thompson, “On the Disparity Between Swish and GELU - Towards Data Science,” Medium, Mar. 03, 2021. <https://towardsdatascience.com/on-the-disparity-between-swish-and-gelu-1ddde902d64b>.